STELLAR SMART CONTRACT AUDITS

# stellar.chat

## AI security audits for Stellar contracts.

Arsenii Pechenkin

arspechen@gmail.com

PROBLEM

# A complement to the manual audit, not a replacement.

Novel economic design and cross-protocol invariants still belong to a human auditor. An AI pass does not take that job. It sits earlier, inside CI/CD and QA, so every change gets a check before a senior reviewer spends days on it.

**$40–100K**
Typical manual DeFi audit

**3–38 days**
Manual audit turnaround

**10–30 min**
Automated pass in CI/CD

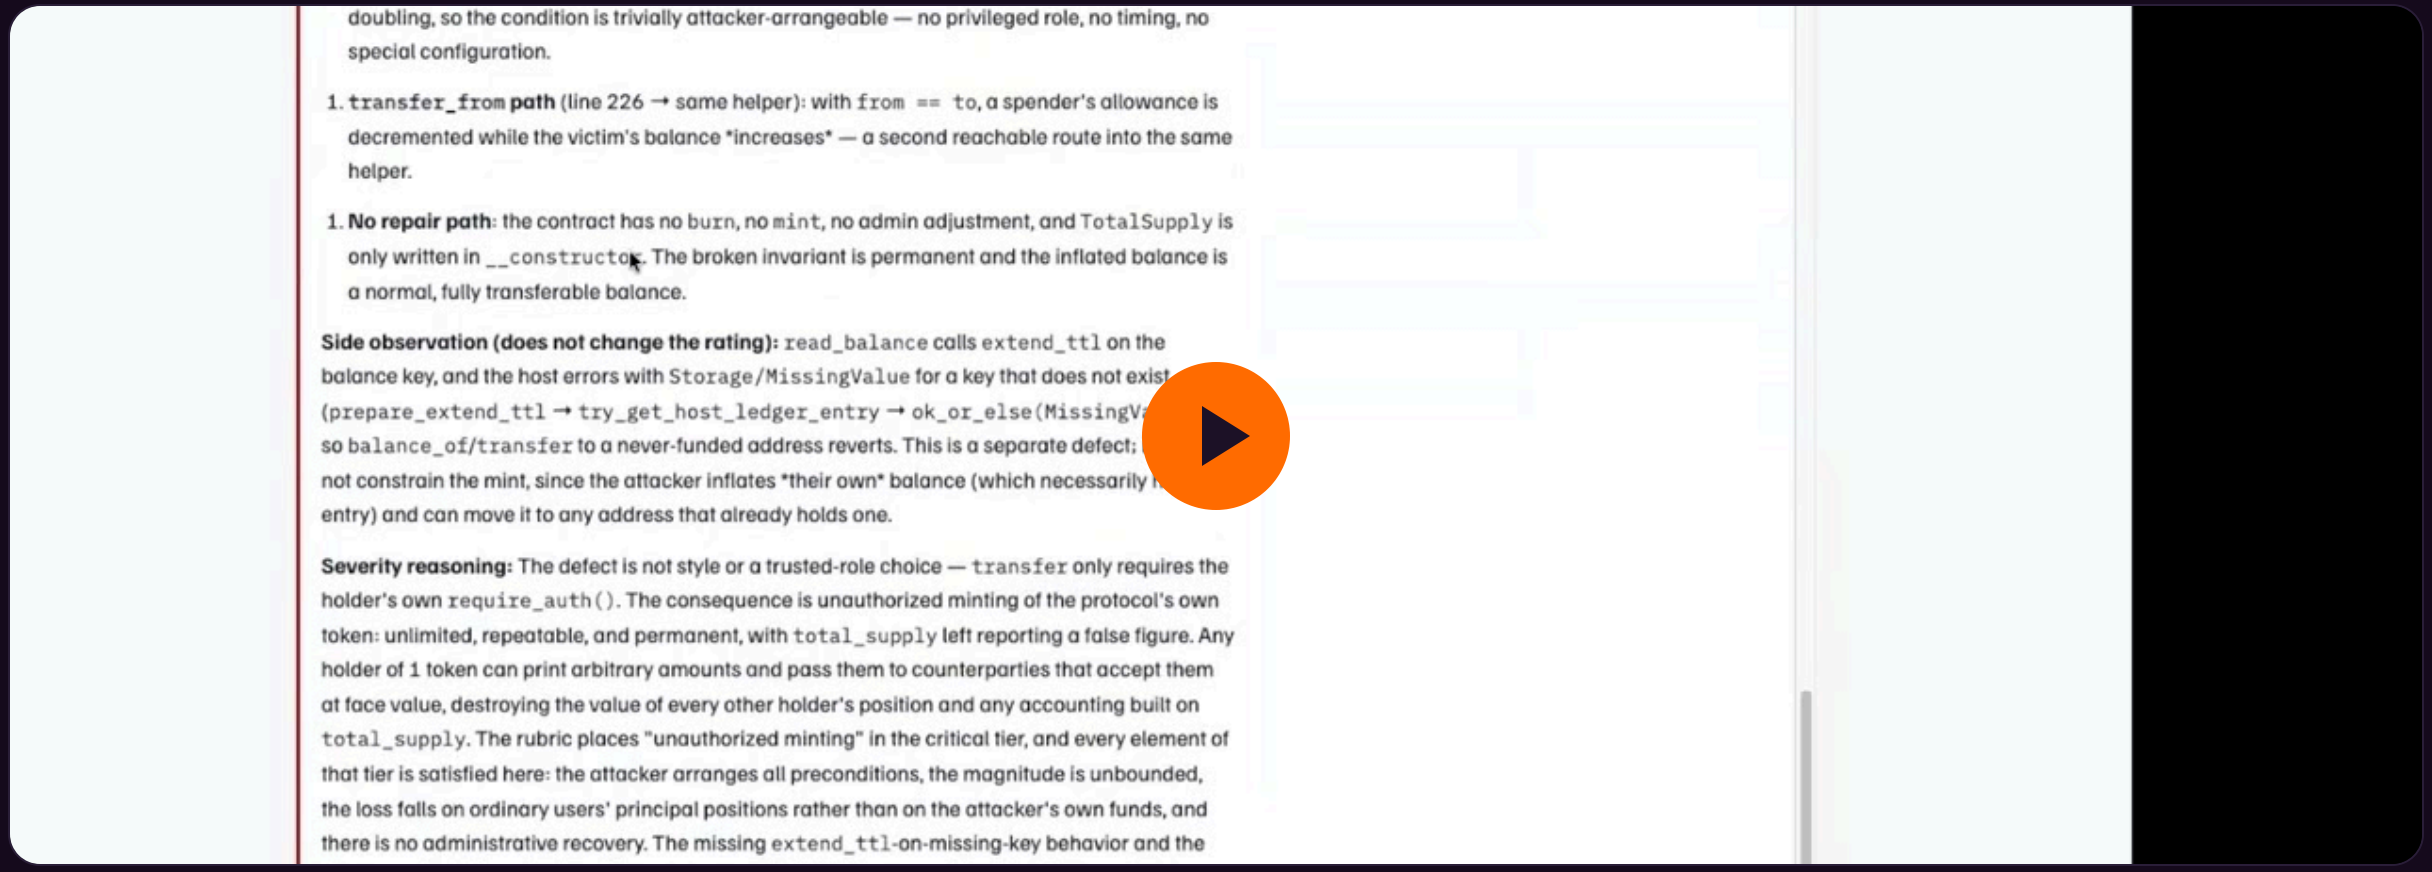
Source: savant.chat, “Faster than a manual audit” and the comparison table. Manual cost is the Sherlock range they cite. These are published ranges, not a quote for this repo.

WHERE IT FITS

Run it on each change in CI/CD and as a QA gate. Keep the human audit for the release. The service filters known bug classes so the auditor’s time goes to what automation still cannot see.

SOLUTION

# Demo




Demo: youtu.be/Me9n0rxoPEw

REPOSITORY

# Source

## github.com/Pe4enable/stellar.chat

Vite and React on the client. A same-origin proxy in the dev server calls Savant, so the API key stays on the machine that launches the app and never reaches the browser.

| IN THE REPO | WHAT IT IS |
|---|---|
| `src/` and `server/audit.mjs` | Editor, tier picker, report, and the Savant REST proxy |
| `contracts/self-transfer-mint` | Soroban token whose self-transfer mints |
| `contracts/fr-nullifier` | Nullifier that accepts an unreduced BN254 scalar |
| `README.md` | Setup: Node, SAVANT_API_KEY, Rust 1.91, soroban-sdk |

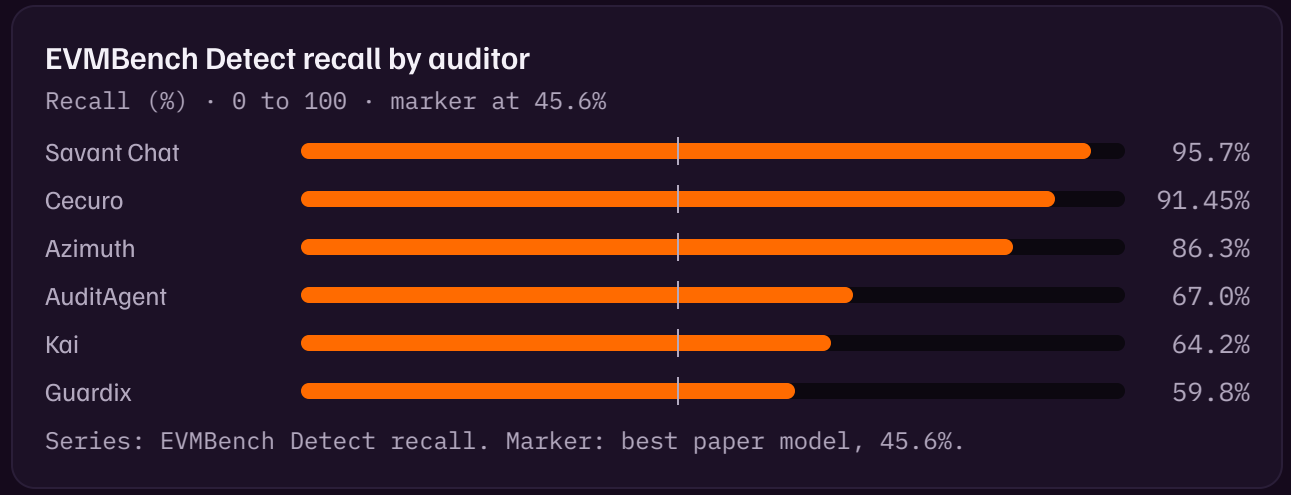
RUN IT

`npm install`, export SAVANT_API_KEY, then `npm run dev`. The app listens on localhost:5173.

BENCHMARKS

# EVMBench Detect recall

EVMBench is the smart-contract security benchmark from OpenAI and Paradigm: 40 Code4rena contests, 117 high-severity bugs. Detect mode scores whether an audit report names the same bug. These rows are published AI auditors, not a score of the stellar.chat UI. stellar.chat calls Savant Chat, which leads the list.




| RANK | AUDITOR | RECALL | FOUND | TYPE |
|---|---|---|---|---|
| 1 | Savant Chat | 95.7% | 112/117 | Agent |
| 2 | Cecuro | 91.45% | — | Agent |
| 3 | Azimuth (TestMachine) | 86.3% | 101/117 | Agent |
| 4 | AuditAgent (Nethermind) | 67.0% | 80/120 | Agent |
| 5 | Kai (Dria) | 64.2% | 77/120 | Agent |
| 6 | Guardix | 59.8% | 70/117 | Agent |

EVMBench Detect recall (%) by published AI auditor. Marker: Claude Opus 4.6 in the EVMBench paper, 45.6%, the best frontier-model result they report. Source: savant.chat/evmbench, leaderboard updated October 20, 2026. Agent scores are vendor-reported except Savant Chat, whose 117 reports and verdicts are public. AuditAgent (80/120) and Kai (77/120) predate OpenAI dropping three findings; the rest use 117.